

# Class 7: Machine Learning 1


Amanda Wilpitz | A17463962

Today we will explore unsupervised machine learning methods including clustering and dimensionality reduction methods.

Let's start by making up some data (where we know there are clear groups that we can use to test out different clustering methods).

We can use the `rnorm()` function to help us here, which randomly generates data with normal distribution of mean and sd.

```
hist(rnorm(n=3000, mean = 3))
```



Now, make data `z` with two “clusters”.


```
x <- c(rnorm(30, mean = -3),
      rnorm(30, mean = +3))

z <- cbind(x=x, y=rev(x))

head(z)
```

```
      x      y
[1,] -1.875590 5.095042
[2,] -3.135288 1.530411
[3,] -2.819450 3.349256
[4,] -3.679415 3.500435
[5,] -3.376655 4.119039
[6,] -2.853156 4.160412
```

```
plot(z)
```



## K-means Clustering

The main function in “base” R for K-means clustering is called `kmeans()`. This will group an input data into input number of clusters that each center around its mean points.

k

```
k$cluster
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2  
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```


Q. Center of each cluster?

```
k$center
```

```
      x      y  
1 -2.946903  2.890691  
2  2.890691 -2.946903
```


Q. Put this result info together and make a little “base R” plot of our clustering result. Also add the cluster center points to this plot.

```
plot(z, col=c("blue", "red"))
```




You can also color by number (1=black):

```
plot(z, col=c(1,2))
```



Color by membership/cluster:

```
plot(z, col=k$cluster)
points(k$centers, col= "blue", pch=16)
```



Q. Run k-means on our input `z` and define 4 clusters making the same results visualization plot as above (plot of `z` colored by cluster membership)

```
k4 <- kmeans(z, centers=4)
k4
```

K-means clustering with 4 clusters of sizes 14, 12, 16, 18

Cluster means:

	x	y
1	2.157563	-3.418860
2	-2.893370	3.871340
3	3.532177	-2.533941
4	-2.982592	2.236925

Clustering vector:

```
[1] 2 4 2 2 2 2 2 4 4 4 2 2 4 4 2 2 4 4 4 4 4 4 2 4 4 4 4 3 1 3 3 3 1 1 1
[39] 1 3 1 3 3 1 1 3 1 3 3 3 1 1 1 1 3 3 3 3 1 3
```


Within cluster sum of squares by cluster:

```
[1] 11.901815  9.782055 15.241262 18.025997
(between_SS / total_SS = 95.1 %)
```

Available components:

[1]	"cluster"	"centers"	"totss"	"withinss"	"tot.withinss"
[6]	"betweenss"	"size"	"iter"	"ifault"	

```
plot(z, col=k$cluster)
points(k$centers, col = "blue", pch=16)
```



Better plot clustering will have a smaller tot.withinss value

```
k$tot.withinss
```

```
[1] 94.19764
```

```
k4$tot.withinss
```

```
[1] 54.95113
```

You can also plot a scree plot (x=number of clusters, y=tot.withinss) and “elbow” in the plot will determine the number of clusters that you should have.

## Hierarchical Clustering

The main function in base R for this is called `hclust()`. It will take as input a distance matrix (you cannot just give your “raw” data as input - you have to first calculate a distance matrix from your data).

```
d <- dist(z)
hc <- hclust(d)
hc
```

Call:

```
hclust(d = d)
```

```
Cluster method   : complete
Distance          : euclidean
Number of objects: 60
```

Plot `hclust`, which will produce a hierarchical tree of the input values. All lower numbers <30 are on the left side and higher numbers 31-60 are on the right side. Each point starts as it's own cluster and the closest points are grouped together until there is one cluster left as the highest branch in the hierarchy. Higher branches = points farther apart.

```
plot(hc)
abline(h=10, col="red")
```



## Cluster Dendrogram



```
hclust (*, "complete")
```


Once I inspect the “tree”, I can “cut” the tree at a certain height to yield my groupings or clusters. The function to do this is called `cutree()`.

```
grps <- cutree(hc, h=10)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
```

```
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
plot(z, col=grps)
```



There are 4 methods to determine distance between clusters in `hclust()` - use trial and error.

## Hands on with Principal Component Analysis (PCA): PCA w/ UK Food

Let's examine some silly 17-dimensional data detailing food consumption in the UK (England, Scotland, Wales, and N. Ireland). Are these countries eating habits different or similar to one another?

### Data import

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
x
```

		X England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

7	Fresh_potatoes	720	874	566	1033
8	Fresh_Veg	253	265	171	143
9	Other_Veg	488	570	418	355
10	Processed_potatoes	198	203	220	187
11	Processed_Veg	360	365	337	334
12	Fresh_fruit	1102	1137	957	674
13	Cereals	1472	1582	1462	1494
14	Beverages	57	73	53	47
15	Soft_drinks	1374	1256	1572	1506
16	Alcoholic_drinks	375	475	458	135
17	Confectionery	54	64	62	41

## Checking your Data

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 5
```

```
dim(x)
```

```
[1] 17  5
```

Set `rownames()` to the first column and removes the first column with the -1 column index:

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586

Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Running this code multiple times will continuously remove the first column.

The code below will keep the first column (foods) as the row name, without deleting any other columns after multiple runs.

```
x <- read.csv(url, row.names=1)
head(x)
```


	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

I prefer the one where you keep the first column as the row name, since it won’t delete other columns especially when you are running the code again and again. This makes it more robust than removing the first column and setting it as the row names.

### Spotting major differences and trends


```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

Changing `beside` to `FALSE` puts the bar plots on top of each other.


```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

The label is the y-axis on its corresponding row and the x-axis in its corresponding column. For example, England is on the y-axis across the first row, Wales in on the x-axis on the second column. When two countries are similar, there are points on the diagonal. If the value is higher for the y-axis, the point is above the diagonal. If the value is higher for the x-axis, the points are below the diagonal.

```
pairs(x, col=rainbow(10), pch=16)
```



Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

N. Ireland has the more differences from the other countries of the UK, since the column and row corresponding to N. Ireland has the most points off the diagonal.

Looking at these types of “pairwise plots” can be helpful but it does not scale well and kind of sucks! There must be a better way...

## PCA to the Rescue!

Principal component analysis (PCA) is a well established “multivariate statistical technique” used to reduce the dimensionality of a complex data set to a more manageable number (typically 2D or 3D). In our example here, we have 17 dimensional data for 4 countries. We can thus “imagine” plotting the 4 coordinates representing the 4 countries in 17 dimensional space. If there is any correlation between the observations (the countries), this will be observed in the 17 dimensional space by the correlated points being clustered close together.

The main function for PCA in base R is called `prcomp()`. This function wants the transpose of our input data - ie. the important foods in as columns and the countries as rows.

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Proportion of variance tells you percentage of data that is captured by that particular axis. PC1 has the highest variance. Cumulative proportion gives proportion of data that is captured by all the axis so far (ie. 96.5% captured by 2 axes)

Let's see what is in our PCA result object `pca`:

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"

$class
[1] "prcomp"
```

The `pca$x` result object is where we will focus first as this details how the countries are related to each other in terms of our new “axis” (aka “PCs”, “eigenvectors”, etc.)


```
head(pca$x)
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.


```
# Plot PC1 vs PC2
plot(pca$x[,1], pca$x[,2], col=c("orange", "red", "blue", "green"),
     xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```





Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
plot(pca$x[,1], pca$x[,2], col=c("orange", "red", "blue", "green"),
     xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x), col=c("orange", "red", "blue", "green"))
```



### Digging deeper (variable loadings)


We can look at the so-called PC “loadings” result object to see how the original foods contribute to our new PCs (ie. how the original variables contribute to our new better PC variables). A positive loading is a positive correlation, while a negative loading is a negative correlation between food and a particular PC.

```
pca$rotation[, 1]
```

Cheese	Carcass_meat	Other_meat	Fish
-0.056955380	0.047927628	-0.258916658	-0.084414983
Fats_and_oils	Sugars	Fresh_potatoes	Fresh_Veg
-0.005193623	-0.037620983	0.401402060	-0.151849942
Other_Veg	Processed_potatoes	Processed_Veg	Fresh_fruit
-0.243593729	-0.026886233	-0.036488269	-0.632640898
Cereals	Beverages	Soft_drinks	Alcoholic_drinks
-0.047702858	-0.026187756	0.232244140	-0.463968168
Confectionery			
-0.029650201			


Plot a bar plot representing PC1

```
par(mar=c(10, 3, 0.35, 0))
barplot(pca$rotation[,1], las=2)
```



Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
par(mar=c(10, 3, 0.35, 0))
barplot(pca$rotation[,2], las=2)
```



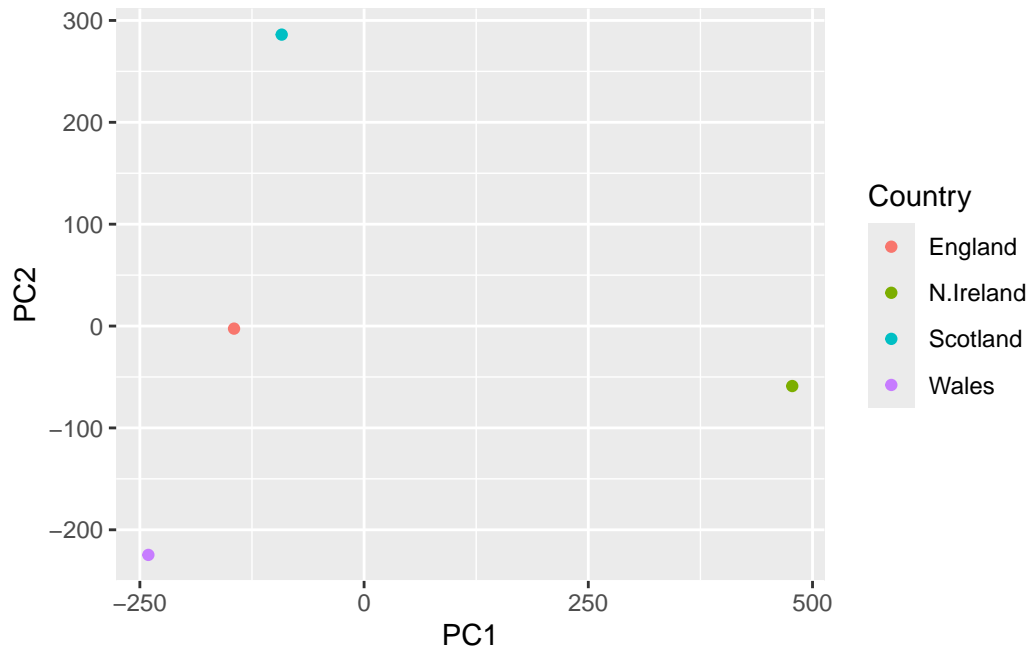
Soft drinks and fresh potatoes predominate. PC2 mainly tells us the contribution of each food category along the PC2 axis, with soft drinks having a positive correlation and fresh potatoes having a negative correlation.

### Using ggplot for these figures

```
library(ggplot2)

df <- as.data.frame(pca$x)
df_lab <- tibble::rownames_to_column(df, "Country")

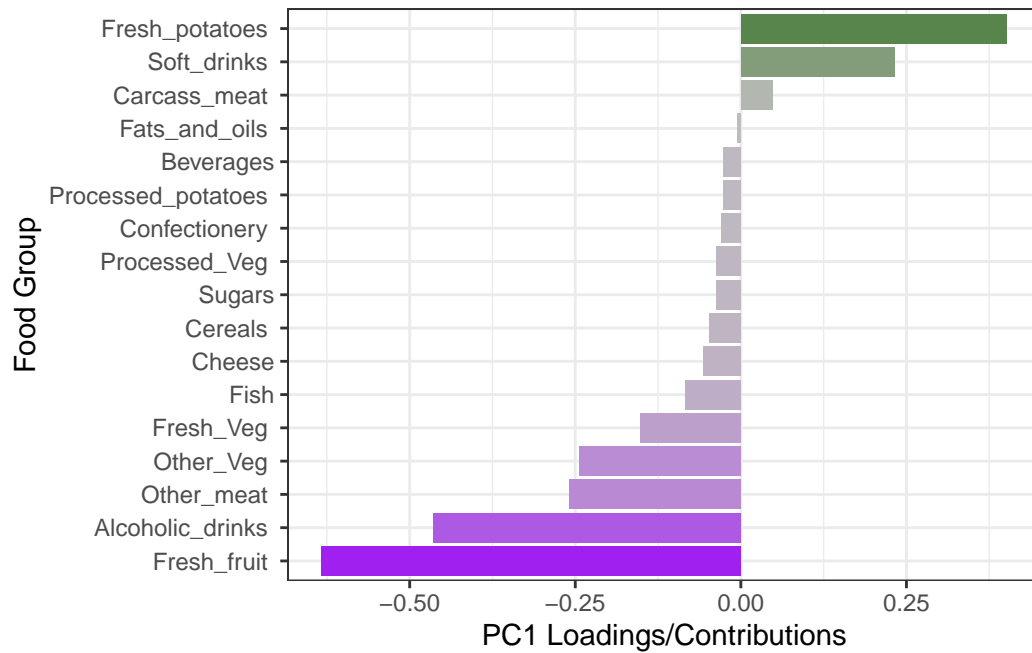
# Our first basic plot
ggplot(df_lab) +
  aes(PC1, PC2, col=Country) +
  geom_point()
```



To plot our loadings plot with ggplot, we will convert it to a data frame and add row names as a new column called “Food”:

```
ld <- as.data.frame(pca$rotation)
ld_lab <- tibble::rownames_to_column(ld, "Food")

ggplot(ld_lab) +
  aes(PC1, reorder(Food, PC1), bg=PC1) +
  geom_col() +
  xlab("PC1 Loadings/Contributions") +
  ylab("Food Group") +
  scale_fill_gradient2(low="purple", mid="gray", high="darkgreen", guide=NULL) +
  theme_bw()
```



## Biplots

Another way to visualize PCA information is in a biplot. The data is organized in a central group of foods around the middle of each PC, with some on the periphery. Points closer to each other are more similar and longer arrows pointing towards/away the PC axis mean the variable contributes more to the PC.


```
biplot(pca)
```



Let's make a PCA and plot the results:

```
pca <- prcomp(t(rna.data), scale=TRUE)

## Simple un polished plot of pc1 and pc2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
```



```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	9.6237	1.5198	1.05787	1.05203	0.88062	0.82545	0.80111
Proportion of Variance	0.9262	0.0231	0.01119	0.01107	0.00775	0.00681	0.00642
Cumulative Proportion	0.9262	0.9493	0.96045	0.97152	0.97928	0.98609	0.99251


  

	PC8	PC9	PC10
Standard deviation	0.62065	0.60342	3.457e-15
Proportion of Variance	0.00385	0.00364	0.000e+00
Cumulative Proportion	0.99636	1.00000	1.000e+00

Quick plot of proportion of variance for each PC:



```
plot(pca, main="Quick scree plot")
```




Let's make our own scree plot:

```
## Variance captured per PC
pca.var <- pca$sdev^2

## Percent variance is often more informative to look at
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)

barplot(pca.var.per, main="Scree Plot",
        names.arg = paste0("PC", 1:10),
        xlab="Principal Component", ylab="Percent Variation")
```



We can see from the summary and the Scree plots that PC1 is where all the action is (92.6%).

Using ggplot to plot our RNA-seq data. Again, we must convert PCA to a data frame first:

```
library(ggplot2)

df <- as.data.frame(pca$x)

df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)

ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)
```

